

Abhimanyu Hans

College Park, MD · ahans1@umiacs.umd.edu · (929) 832-1886 · <https://ahans30.github.io/>

RESEARCH INTERESTS

My research interests revolve around ensuring the security, efficiency and robustness of generative models. I particularly care about memorization and privacy related problems around language models.

EDUCATION

University of Maryland – College Park

MS Computer Science

Fall 2022 – May 2025 (Expected)

- Advised by **Prof. Tom Goldstein**

University of Delhi

BS Mathematics (Honours)

2014 – 2017

PUBLICATION

- Spotting LLMs With Binoculars: Zero-Shot Detection of Machine-Generated Text (*ICML*) (*code*) (*demo*)
A. Hans, A. Schwarzschild, V. Cherepanova, H. Kazemi, A. Saha, M. Goldblum, J. Geiping, T. Goldstein
- Be like a Goldfish, Don't Memorize! Mitigating Memorization in Generative LLMs (*preprint*) (*code*)
A. Hans, Y. Wen, J. Kirchenbauer, H. Kazemi, P. Singhanian, S. Singh, G. Somepalli, J. Geiping, A. Bhatele, T. Goldstein

RELEVANT WORK EXPERIENCE

University of Maryland Institute for Advanced Computer Studies

Graduate Research Assistant

College Park, MD

May 2023 - Present

- Working on language models from security and alignment perspective.

PayPal

Machine Learning Engineer

Remote Work (India)

Oct 2021 - Aug 2022

- Automated model monitoring framework, saving 70+ FTE weeks, and deploying interactive dashboards for multiple model-specific teams (spanning 100+ model use cases)
- Employed quantitative metrics (AUC, Accuracy, KS-statistic, population stability index) and qualitative analysis (LIME) to enhance model performance and feature interpretability.

HDFC Bank

Machine Learning Engineer

Mumbai, India

Jan 2020 - Oct 2021

- Created an anomalous transaction detection system using AutoEncoder and Isolation Forest algorithms, enhancing fraud detection capabilities to achieve 25% recall at 0.1% FPR
- Developed a user-friendly model repository CRUD application using Flask, Jinja Templates, and PostgreSQL, streamlining internal model management processes.

Think360.ai

Associate Data Scientist

Mumbai, India

Oct 2018 - Jan 2020

- Developed a random Forest model for predicting repossession values of construction equipment, achieving a 0.84 adjusted R-squared score.
- Developed an automated business intelligence pipeline using Tableau, R, and Bash, serving over 5000 FTE with dashboards and email reports.

SKILLS

- Languages & Tools:** Python, Java, JavaScript, SQL, Git, Docker, Linux, AWS, Shell
- Frameworks:** PyTorch, HuggingFace suite, Flask, React, Tableau, Pandas, NumPy, Sklearn, SpringMVC

PROJECTS (MACHINE LEARNING & SOFTWARE ENGINEERING)

Multi-Armed Bandit Algorithms

(*code*)

A *NumPy*-based implementation of Greedy, Epsilon-Greedy, and Upper Confidence Bound algorithms along with their empirical performance analysis.

Active Learning

(*code*) (*slides*)

An approach to ML wherein the learner queries the user with best instance for a label.

Model Repository

(*code*)

Full-stack Flask app serving as a repository for models and their performance metadata.